

# Distillation Unleashed: Domain Knowledge Transfer with Compact Neural Networks

Hadi Abdi Khojasteh

hadi.abdikhojasteh@**deltatre**.com


September 16th, 2023




Presentation will be available online.




# Overview



Introduction





Variations




Implementation


# Introduction




# Deep Neural Network Demonstration




# Transfer Learning




# Transfer Learning




# Overview



Introduction





Variations





Implementation

# Knowledge Transfer

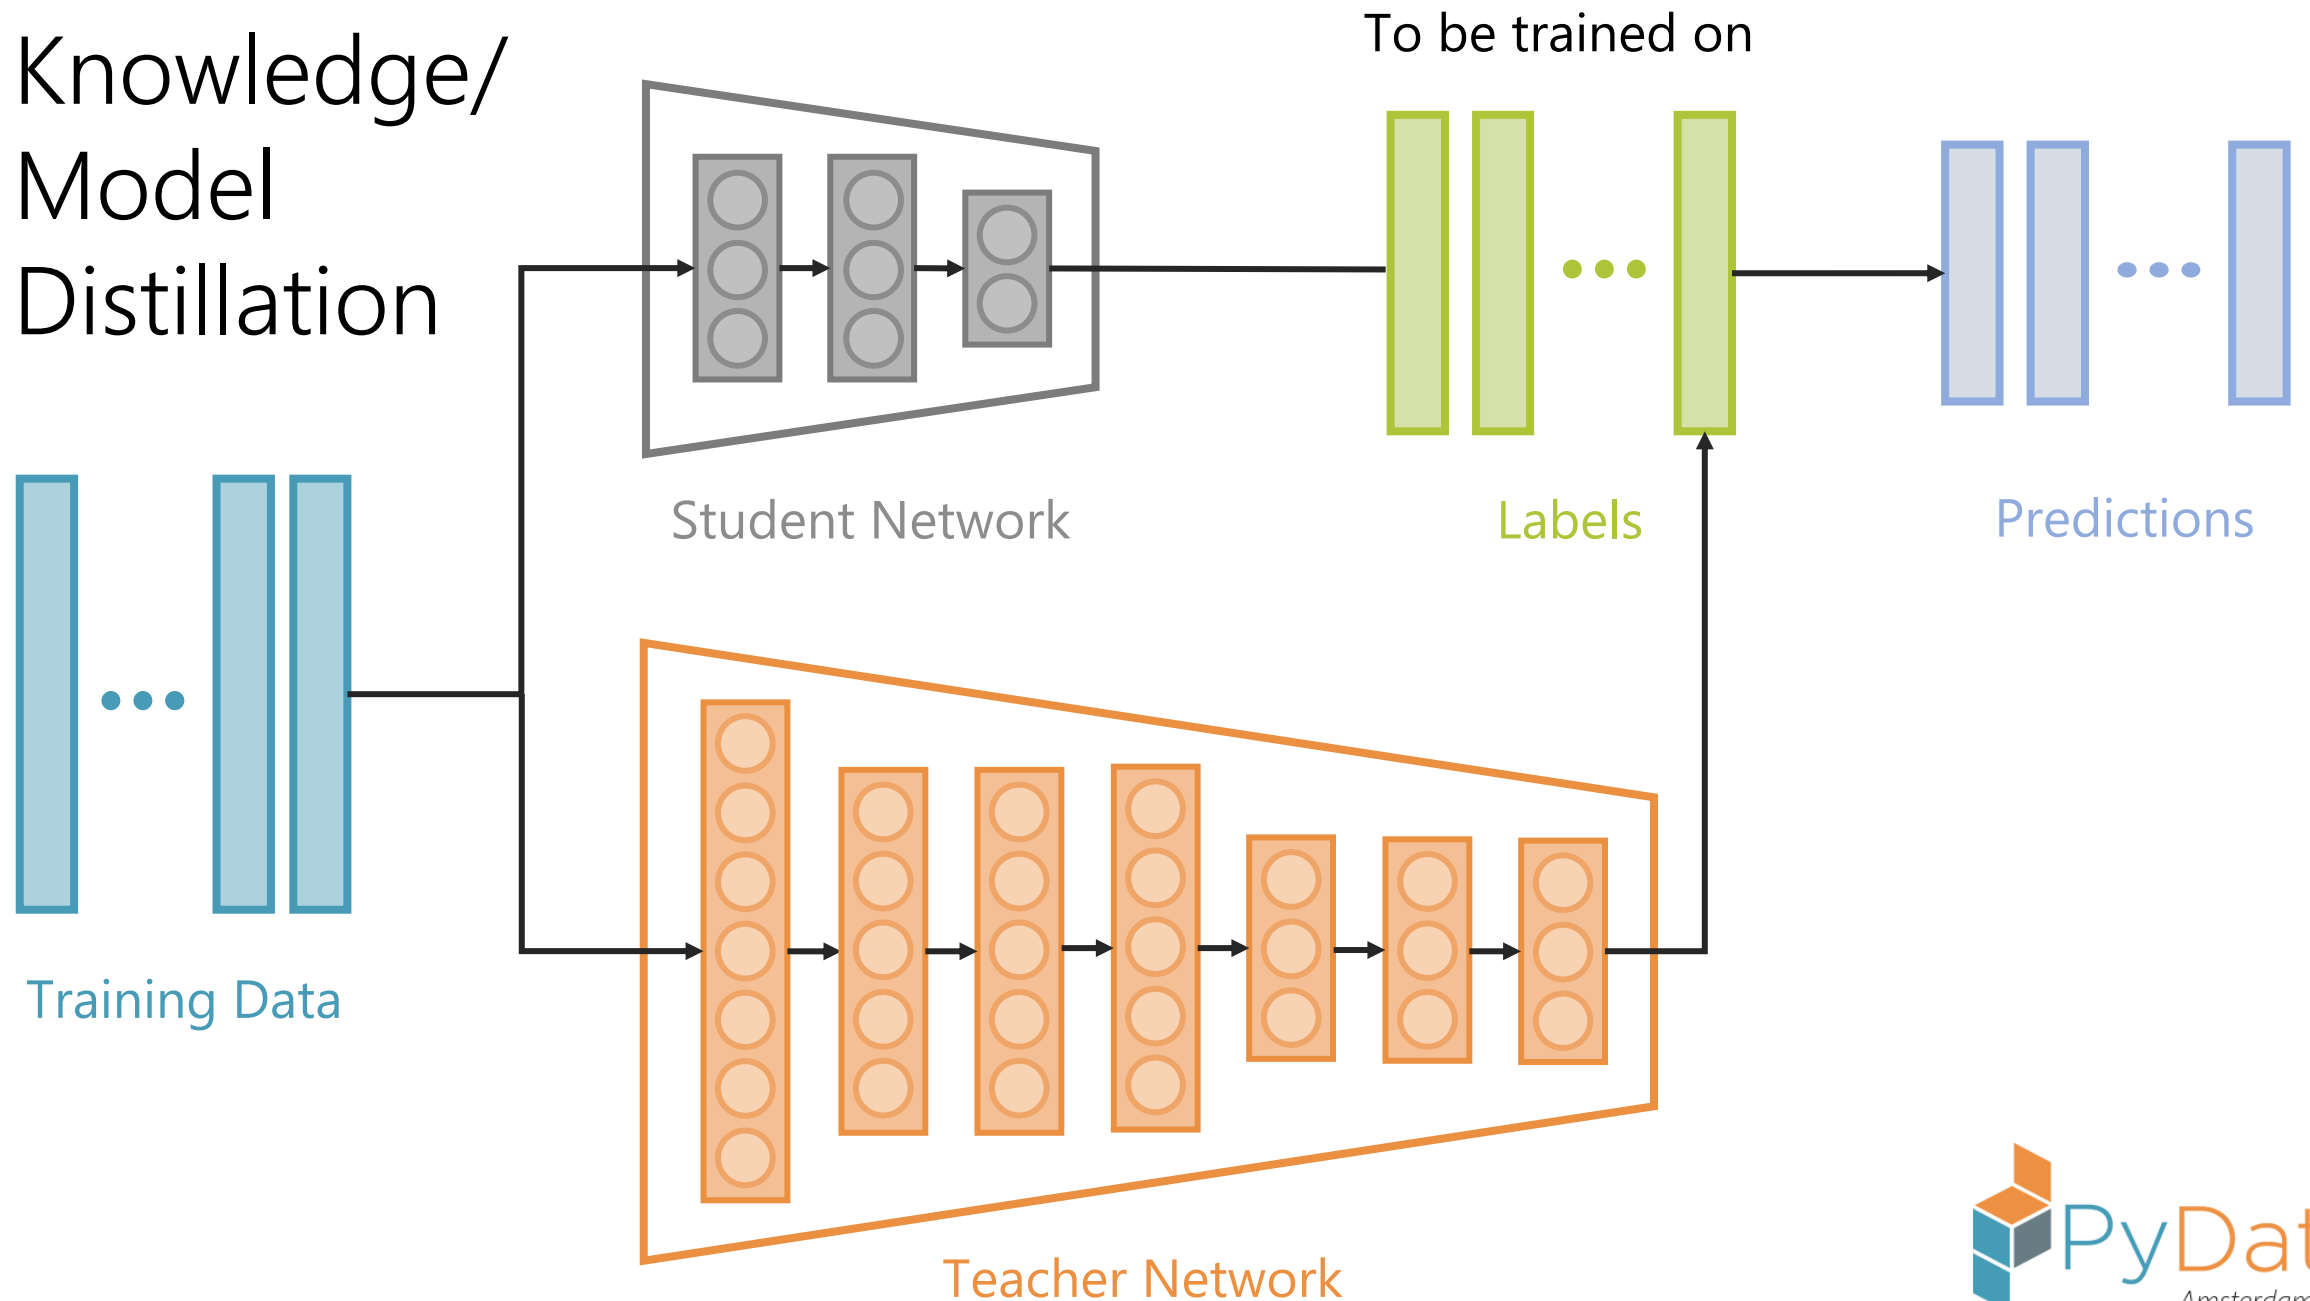




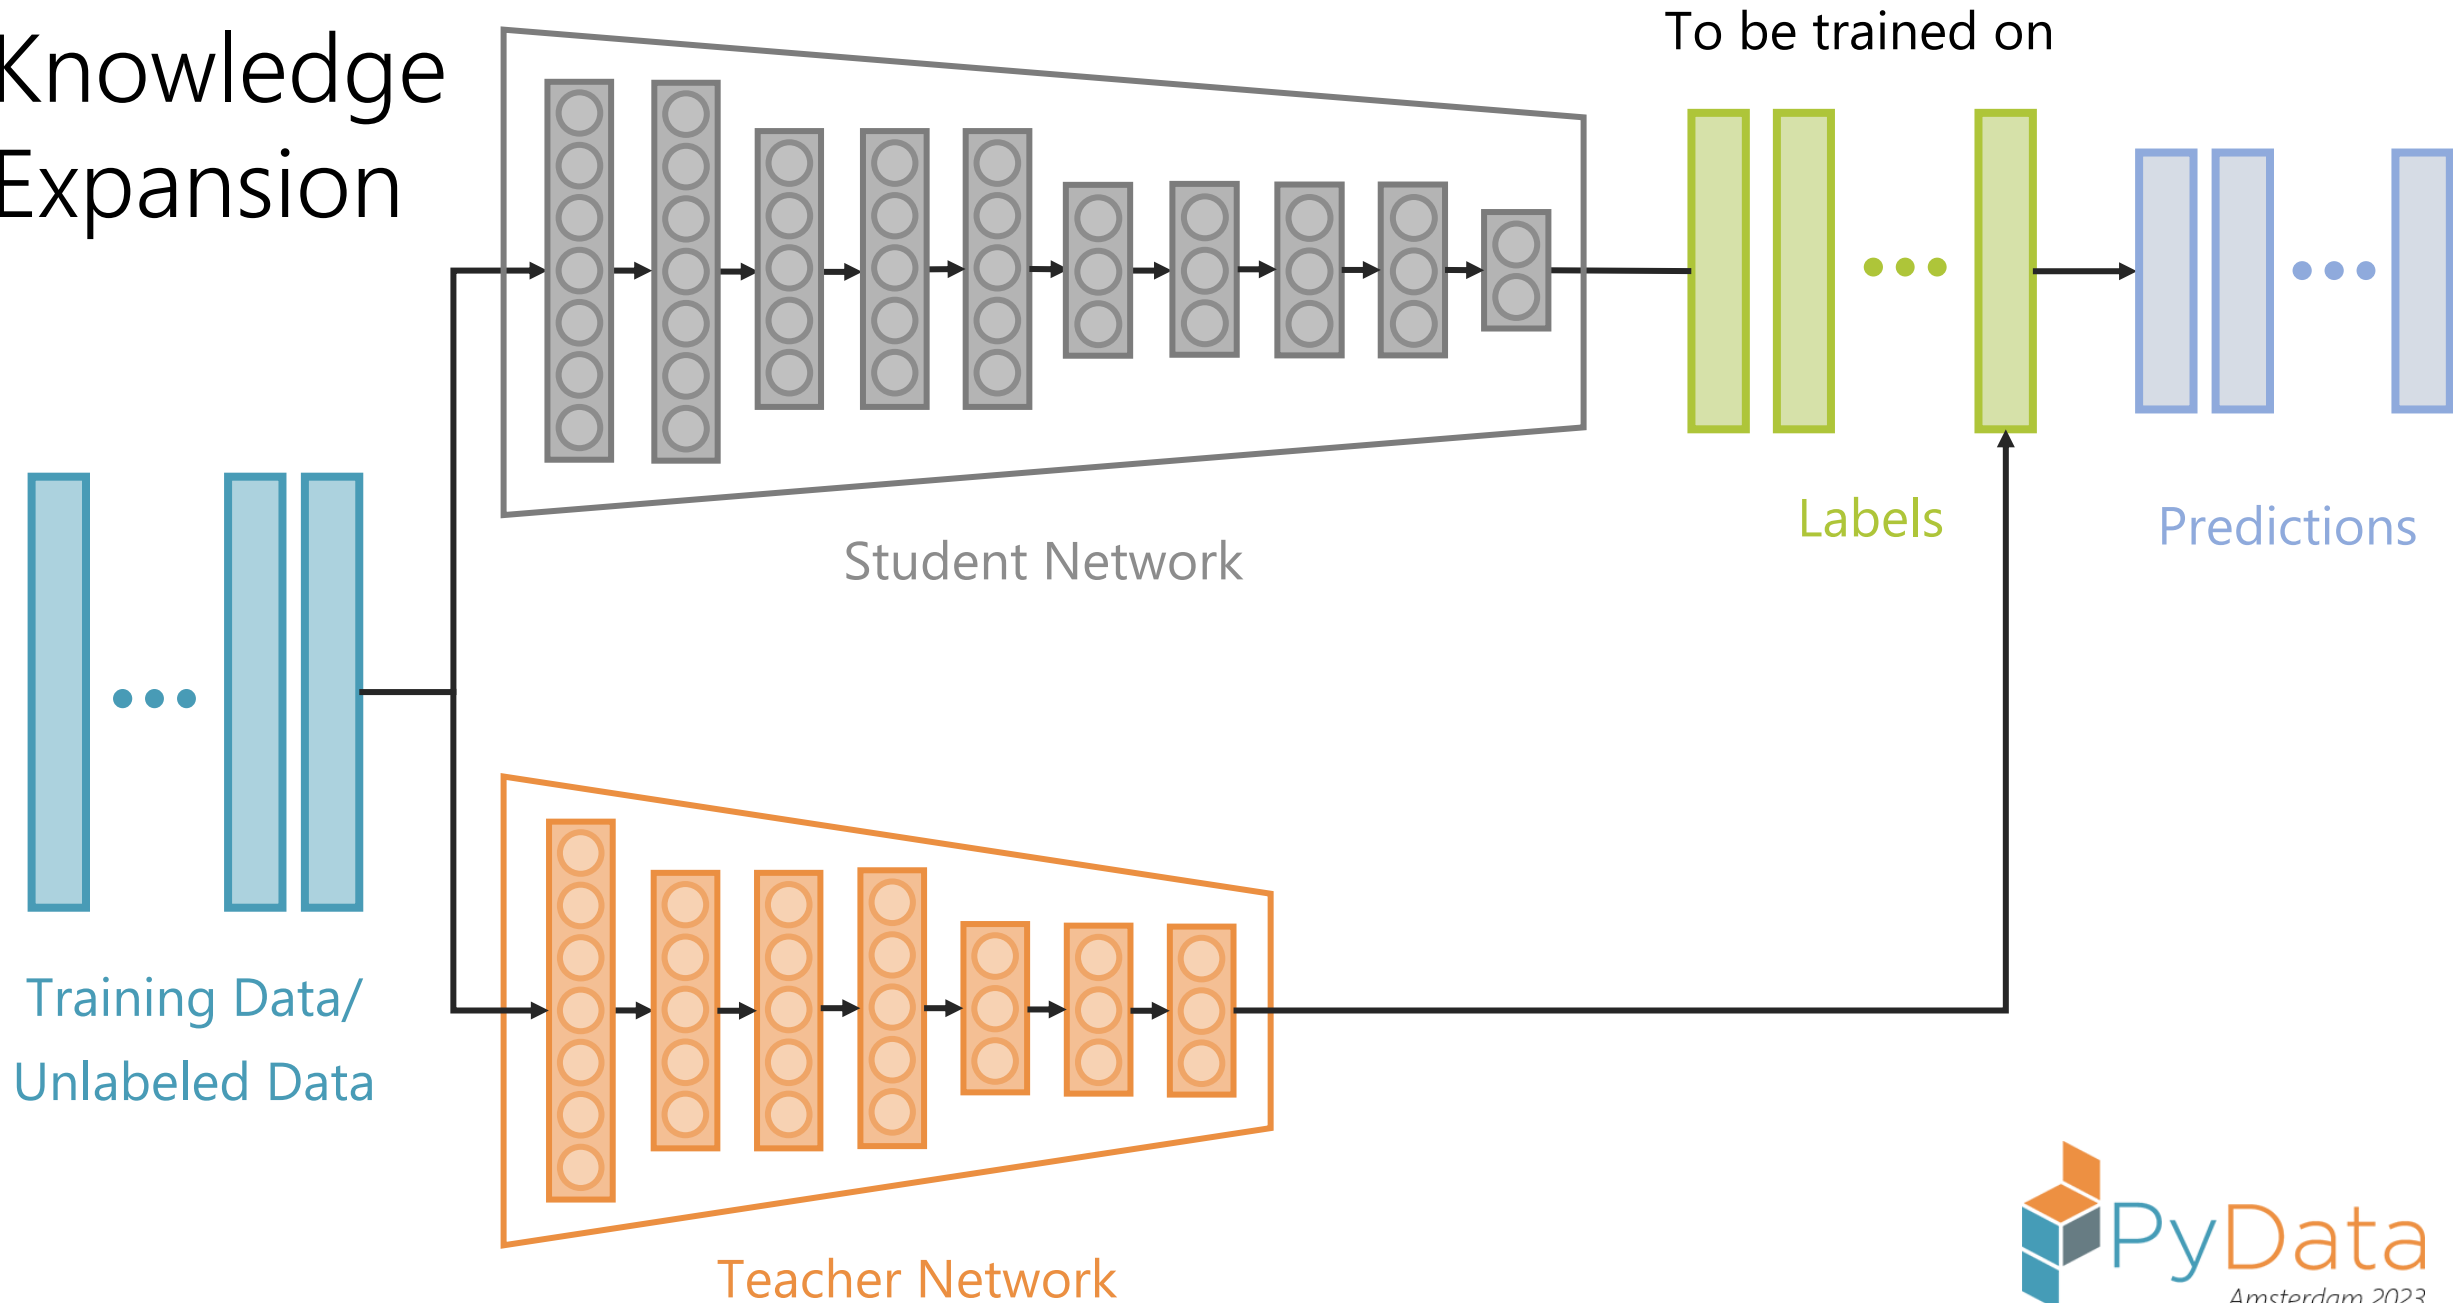
# Knowledge Transfer



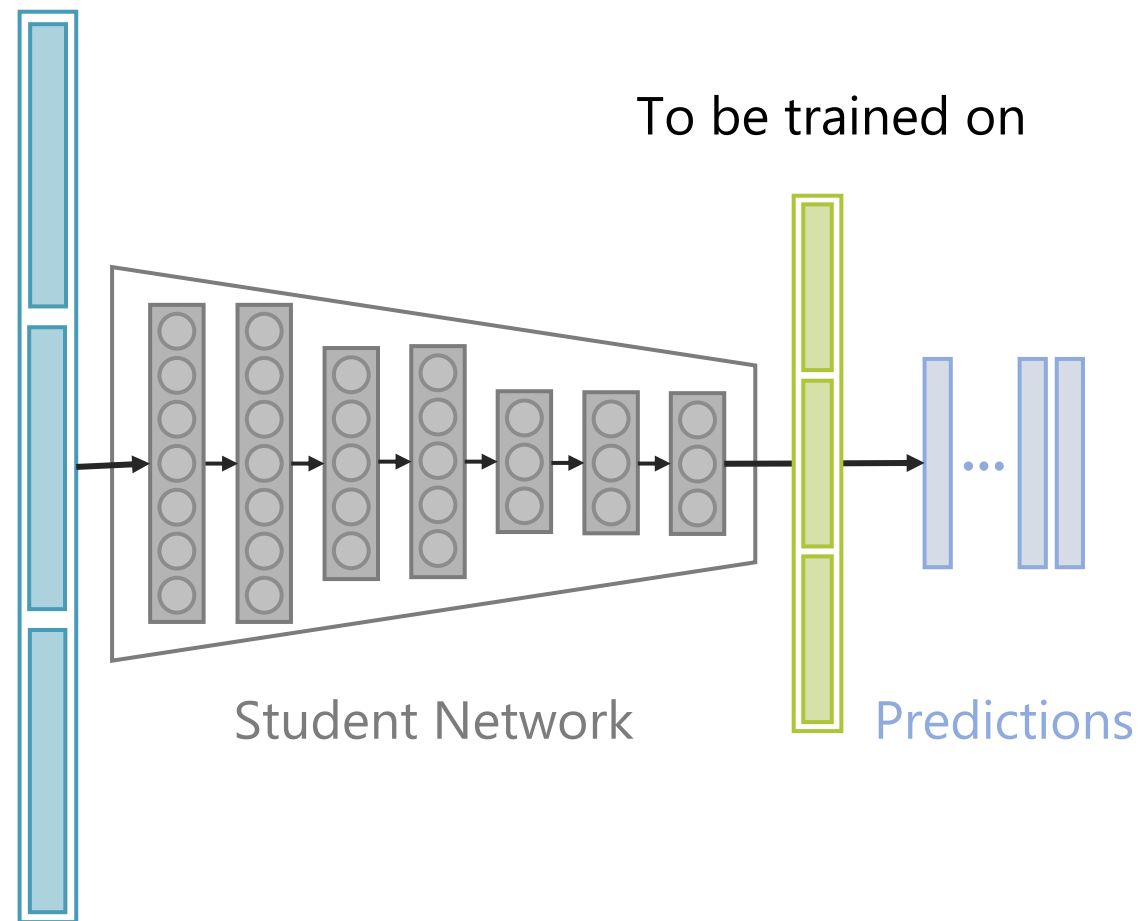

# Knowledge/ Model Distillation




# Knowledge Expansion




# Multi-task Distillation




# Overview



Introduction



Variations




Implementation

# Distillation Learning in Action

First, we import the TensorFlow library:


```
# Importing libraries
import tensorflow as tf
```



Then load MNIST dataset and normalize the data:

```
# Loading the data
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
```

```
# Preprocessing the data
x_train = x_train.reshape(-1, 784).astype('float32') / 255
x_test = x_test.reshape(-1, 784).astype('float32') / 255
y_train = tf.keras.utils.to_categorical(y_train)
y_test = tf.keras.utils.to_categorical(y_test)
```



0,0,0,0,0,0,0,0,0,1  
0 1 2 3 4 5 6 7 8 9

Next, we define the teacher model:

# Defining the teacher model (More complex and larger neural network)

```
class TeacherModel(tf.keras.Model):
```

```
    def __init__(self):
```

```
        super(TeacherModel, self).__init__()
```

```
        self.layer1 = tf.keras.layers.Dense(512, input_shape=(784,), activation='relu')
```

```
        self.layer2 = tf.keras.layers.Dense(256, activation='relu')
```

```
        self.layer3 = tf.keras.layers.Dense(128, activation='relu')
```

```
        self.layer4 = tf.keras.layers.Dense(64, activation='relu')
```

```
        self.last = tf.keras.layers.Dense(10, activation='softmax')
```

```
    def call(self, x):
```

```
        # Defining the forward pass
```


```
        x = self.layer1(x)
```

```
        x = self.layer2(x)
```

```
        x = self.layer3(x)
```

```
        x = self.layer4(x)
```

```
        return self.last(x)
```




Teacher Network

And an student model which is shallower than teacher model:

# Defining the student model (More efficient and smaller neural network)

```
class StudentModel(tf.keras.Model):  
    def __init__(self):  
        super(StudentModel, self).__init__()  
        self.layer1 = tf.keras.layers.Dense(512, input_shape=(784,), activation='relu')  
        self.layer2 = tf.keras.layers.Dense(256, activation='relu')  
        self.last = tf.keras.layers.Dense(10, activation='softmax')  
  
    def call(self, x):  
        # Defining the forward pass  
        x = self.layer1(x)  
        x = self.layer2(x)  
        return self.last(x)
```



Student Network



Afterward we call and load the teacher model first, define some of its hyperparameters and finally train it as follows:

```
# Loading the teacher model
```

```
teacher_model = TeacherModel()
```

```
# Defining the loss function and optimizer
```


```
loss_fn = tf.keras.losses.CategoricalCrossentropy()
```

```
optimizer = tf.keras.optimizers.Adam()
```

```
# Training the teacher model
```

```
teacher_model.compile(optimizer=optimizer, loss=loss_fn, metrics=['accuracy'])
```

```
teacher_model.fit(x_train, y_train, epochs=5, batch_size=32,\n                  validation_data=(x_test, y_test))
```




```
Epoch 1/5 1875/1875 - loss: 0.2074 - accuracy: 0.9377 - val_loss: 0.1059 - val_accuracy: 0.9694
Epoch 2/5 1875/1875 - loss: 0.0941 - accuracy: 0.9717 - val_loss: 0.0878 - val_accuracy: 0.9736
Epoch 3/5 1875/1875 - loss: 0.0675 - accuracy: 0.9795 - val_loss: 0.0722 - val_accuracy: 0.9773
Epoch 4/5 1875/1875 - loss: 0.0530 - accuracy: 0.9842 - val_loss: 0.0808 - val_accuracy: 0.9776
Epoch 5/5 1875/1875 - loss: 0.0445 - accuracy: 0.9864 - val_loss: 0.0813 - val_accuracy: 0.9782
```

When teacher model training is done, we freeze teacher model layers to be able to use it in the student model:

```
# Loading the student model
student_model = StudentModel()

# Freezing the teacher model layers
for layer in teacher_model.layers:
    layer.trainable = False
```



The loss function for student model is calculated with respect to the teacher model output not the true values from the training data:

```
# Defining the distillation loss function
```

```
temp = 5
```

```
def distillation_loss(y_true, y_pred):
```

```
    y_true = tf.nn.softmax(y_true / temp)
```


```
    y_pred = tf.nn.softmax(y_pred / temp)
```

```
    return tf.reduce_mean(tf.keras.losses.categorical_crossentropy(y_true, y_pred))
```

```
# Train the student model
```

```
student_model.compile(optimizer=optimizer, loss=distillation_loss, metrics=['accuracy'])
```

```
student_model.fit(x_train, teacher_model.predict(x_train), epochs=5, batch_size=32,  
validation_data=(x_test, y_test))
```



```

1875/1875
Epoch 1/5 - loss: 2.3008 - accuracy: 0.9404 - val_loss: 2.3007 - val_accuracy: 0.9643
Epoch 2/5 - loss: 2.3007 - accuracy: 0.9741 - val_loss: 2.3007 - val_accuracy: 0.9721
Epoch 3/5 - loss: 2.3007 - accuracy: 0.9796 - val_loss: 2.3007 - val_accuracy: 0.9719
Epoch 4/5 - loss: 2.3007 - accuracy: 0.9837 - val_loss: 2.3007 - val_accuracy: 0.9761
Epoch 5/5 - loss: 2.3007 - accuracy: 0.9862 - val_loss: 2.3007 - val_accuracy: 0.9760
      #loss: 0.0445 - accuracy: 0.9864 - val_loss: 0.0813 - val_accuracy: 0.9782

```

Finally, we run the inference on the dataset:

```


# Evaluate the student model
test_loss, test_acc = student_model.evaluate(x_test, y_test)
print('Test accuracy:', test_acc)

```


```

313/313 [=====] - 1s 2ms/step - loss: 2.3007 - accuracy: 0.9760
Test accuracy: 0.9760000109672546


```




# Overview



Introduction



Variations



Implementation

# Further References

Below are the recommended researches for a comprehensive exploration and in-depth understanding of the topic:

KD++	Improving Knowledge Distillation via Regularizing Feature Norm and Direction, 2023
Closer Look	A closer look at the training dynamics of knowledge distillation, 2023
DIST	Knowledge Distillation from A Stronger Teacher, 2022
ADLIK-Faster	Focal and Global Knowledge Distillation for Detectors, 2021
LSHFM	Distilling Knowledge by Mimicking Features, 2020
ADLIK-MO-P25	Ensemble Knowledge Distillation for Learning Improved and Efficient Networks, 2019
KD/ADLIK-MO	Distilling the Knowledge in a Neural Network, 2015

# Further References

**torchdistill** (formerly kdkit) simplifies knowledge distillation through declarative config files, eliminating the need for extensive code. It also supports reproducible deep learning studies and offers flexibility for various experiments.

**# Fine-tune Transformer models for GLUE CoLA task**

```
!accelerate launch torchdistill/examples/legacy/hf_transformers/text_classification.py \  
  --config torchdistill/configs/legacy/sample/glue/cola/ce/bert_base_uncased.yaml \  
  --task cola \  
  --log log/glue/cola/ce/bert_base_uncased.txt \  
  --private_output leaderboard/glue/standard/bert_base_uncased/
```

# Thank you for your attention :)

Distillation Unleashed:  
Domain Knowledge Transfer with Compact Neural Networks

You can find me in   



Presentation will be available online.







# Transfer Learning

